

## CSS533: Program 2

Shreevatsa Ganapathy Hegde

University of Washington, Bothell

[sghegde@uw.edu](mailto:sghegde@uw.edu)

Prof. Munehiro Fukuda

*1st May, 2025*

## Table of Contents

1	Documentation
2	Source Code
3	Execution
4	Discussion
5	Lab- 4

## Documentation:

The Documentation section involves the implementation details of the Program 2: Flight Data Analysis with Apache Storm that I have built. This section contains the implementation of the requirements that were given in the program 2 document. The additional features will be discussed in the discussion section.

*TopologyMain Class:* This class sets up and executes a Storm topology locally for processing flight data. It uses a spout (FlightsDataReader) and connects it to two bolts (HubIdentifier and AirlineSorter) to group and analyze flights based on proximity to airports and airline information. The main() method accepts two arguments: the first is the path to flights.txt with flight data, and the second is the path to airports.txt containing airport information. The topology runs on a local Storm cluster for 10 seconds before shutting down.

*FlightsDataReader Class:* This class is a custom Storm spout that reads a JSON-formatted flight dataset from a file and emits each flight record as a tuple to the topology. In the open() method, it initializes the file reader using the path provided in the Storm configuration. The nextTuple() method reads the entire file once, parses the JSON, and iterates over the "states" array, which contains individual flight entries. For each flight with 17 fields, it emits a tuple with those fields (handling nulls as empty strings) using the SpoutOutputCollector. The declareOutputFields() method defines the schema of each emitted tuple (e.g., call sign, latitude, altitude, etc.). The class also implements ack(), fail(), and close() to handle basic message acknowledgement and cleanup.

*HubIdentifier class:* The HubIdentifier class is a Storm bolt that filters flights to identify those near major airports and emits the airport-city, airport-code, and call-sign for each. It loads airport data during setup and, for each flight, checks whether it is within 20 miles in latitude or longitude of any known airport. If a nearby airport is found, the flight is passed on for further processing.

*AirlineSorter class:* This class is a Storm bolt that groups and counts flights by airline code for each airport. It receives tuples containing the airport city, airport code, and flight call sign. For each tuple, it extracts the first three characters of the call sign as the airline code and updates a count of flights for that airline at the corresponding airport. In the cleanup() method, it prints out a summary of the total flights per airline at each airport, along with the total number of flights at that airport. *Counterpart Class:* This thread runs in the background and handles the opponent's actions during the game. It waits for input from the other player, marks their move on the board, and checks for a win or draw. If the game ends, it shows the result, and disables the board. It also flips the turn back to the player after each valid move.

A general data flow goes as this. The topology processes flight data by first reading a JSON-formatted flights.txt file through the FlightsDataReader spout, which emits tuples containing flight details. These are passed to the HubIdentifier bolt, which loads airport location data from airports.txt and filters flights to identify those within 20 miles (in latitude or longitude) of a listed airport. For each qualifying flight, it emits the nearest airport and the flight's call sign to the AirlineSorter bolt, which groups flights by airport and airline code, tracks counts, and prints a summary of airline activity per airport when the topology finishes.

## Source Code:

### TopologyMain Class:

```
import backtype.storm.Config;
import backtype.storm.LocalCluster;
import backtype.storm.topology.TopologyBuilder;
import backtype.storm.tuple.Fields;

import spouts.FlightsDataReader;
import bolts.HubIdentifier;
import bolts.AirlineSorter;

public class TopologyMain {
    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Usage: TopologyMain <flights.txt path> <airports.txt path>");
            System.exit(1);
        }

        String flightsPath = args[0];
        String airportsPath = args[1];

        // Create topology
        TopologyBuilder builder = new TopologyBuilder();

        builder.setSpout("flights-reader", new FlightsDataReader());

        builder.setBolt("hub-identifier", new HubIdentifier())
            .shuffleGrouping("flights-reader");

        builder.setBolt("airline-sorter", new AirlineSorter(), 1)
            .fieldsGrouping("hub-identifier", new Fields("airport.city"));

        // Configuration
        Config conf = new Config();
        conf.put("FlightsFile", flightsPath);
        conf.put("AirportsData", airportsPath);
        conf.setDebug(false);
        conf.put(Config.TOPOLOGY_MAX_SPOUT_PENDING, 1);

        // Run locally
        LocalCluster cluster = new LocalCluster();
        cluster.submitTopology("Flight-Data-Analysis", conf, builder.createTopology());

        Thread.sleep(10000);
        cluster.shutdown();
    }
}
```

### FlightsDataReader Class:

```
package spouts;

import backtype.storm.spout.SpoutOutputCollector;
import backtype.storm.task.TopologyContext;
import backtype.storm.topology.OutputFieldsDeclarer;
```

```

import backtype.storm.topology.base.BaseRichSpout;
import backtype.storm.tuple.Fields;
import backtype.storm.tuple.Values;

import java.io.BufferedReader;
import java.io.FileReader;
import java.util.Map;
import org.json.JSONArray;
import org.json.JSONObject;

public class FlightsDataReader extends BaseRichSpout {
    private SpoutOutputCollector collector;
    private FileReader fileReader;
    private boolean completed = false;

    public void ack(Object msgId) {
        System.out.println("OK: " + msgId);
    }

    public void fail(Object msgId) {
        System.out.println("FAIL: " + msgId);
    }

    public void close() {}

    public void nextTuple() {
        if (completed) {
            try {
                Thread.sleep(1000);
            } catch (InterruptedException e) {}
            return;
        }

        try {
            StringBuilder jsonBuilder = new StringBuilder();
            BufferedReader reader = new BufferedReader(fileReader);
            String line;
            while ((line = reader.readLine()) != null) {
                jsonBuilder.append(line);
            }

            JSONObject root = new JSONObject(jsonBuilder.toString());
            JSONArray states = root.getJSONArray("states");

            for (int i = 0; i < states.length(); i++) {
                JSONArray flight = states.getJSONArray(i);
                if (flight.length() == 17) {
                    Object[] fields = new Object[17];
                    for (int j = 0; j < 17; j++) {
                        fields[j] = (flight.isNull(j)) ? "" : flight.get(j).toString().trim();
                    }
                    // System.out.println("EMITTING: " + Arrays.toString(fields));
                    collector.emit(new Values(fields), fields[1]);
                }
            }

        } catch (Exception e) {
            e.printStackTrace();
            throw new RuntimeException("Error parsing flights.json", e);
        }
    }
}

```

```

    }

    completed = true;
}

public void open(Map conf, TopologyContext context, SpoutOutputCollector collector) {
    this.collector = collector;
    try {
        String filePath = conf.get("FlightsFile").toString();
        System.out.println("Reading flights data from: " + filePath);
        this.fileReader = new FileReader(filePath);
    } catch (Exception e) {
        throw new RuntimeException("Error reading flight file", e);
    }
}

public void declareOutputFields(OutputFieldsDeclarer declarer) {
    declarer.declare(new Fields(
        "transponder_address", "call_sign", "origin_country", "last_contact1", "last_contact2",
        "longitude", "latitude", "altitude_barometric", "on_ground", "velocity", "heading",
        "vertical_rate", "sensors", "altitude_geometric", "transponder_code", "special_purpose",
"origin"
    )); // Adjusted couple of names to make them more meaningful and easy to fetch
}
}

```

## HubIdentifier Class:

```

package bolts;

import backtype.storm.topology.BasicOutputCollector;
import backtype.storm.topology.OutputFieldsDeclarer;
import backtype.storm.topology.base.BaseBasicBolt;
import backtype.storm.tuple.Fields;
import backtype.storm.tuple.Tuple;
import backtype.storm.tuple.Values;
import backtype.storm.task.TopologyContext;

import java.io.BufferedReader;
import java.io.FileReader;
import java.io.Serializable;
import java.util.ArrayList;
import java.util.List;
import java.util.Map;

public class HubIdentifier extends BaseBasicBolt {

    private List<Airport> airports = new ArrayList<>();

    public void prepare(Map stormConf, TopologyContext context) { // Initialize the airports list
        String airportsFile = stormConf.get("AirportsData").toString();
        try (BufferedReader br = new BufferedReader(new FileReader(airportsFile))) {
            String line;

```

```

        while ((line = br.readLine()) != null) {
            // There are only 40 entries in the airports.txt file, So there was no need to select top
40 entries.
            String[] parts = line.split(",");
            if (parts.length >= 4) {
                String city = parts[0].trim();
                String code = parts[1].trim();
                double lat = Double.parseDouble(parts[2].trim());
                double lon = Double.parseDouble(parts[3].trim());
                airports.add(new Airport(city, code, lat, lon));
            }
        }
    } catch (Exception e) {
        throw new RuntimeException("Error reading airports file", e);
    }
}

public void execute(Tuple input, BasicOutputCollector collector) {
    try {
        String callSign = input.getStringByField("call_sign").trim();
        double flightLat = Double.parseDouble(input.getStringByField("latitude"));
        double flightLon = Double.parseDouble(input.getStringByField("longitude"));
        double verticalRate = Double.parseDouble(input.getStringByField("vertical_rate"));
        double velocity = Double.parseDouble(input.getStringByField("velocity"));
        double altitude = Double.parseDouble(input.getStringByField("altitude_geometric"));
        boolean onGround = Boolean.parseBoolean(input.getStringByField("on_ground"));

        if (callSign == null || callSign.trim().isEmpty()) {
            return; // Skip flights with no call sign
        }

        // Filter out flyovers
        if (!onGround && verticalRate == 0) {
            return; // Skip flights that are not on the ground and have no vertical rate
        }

        if (velocity > 200) {
            return; // Skip high-speed flights
        }

        if (altitude > 1000) {
            return; // Skip high altitude flights
        }

        Airport nearest = null;
        double minDistance = Double.MAX_VALUE;

        double latThreshold = 20.0 / 70.0; // 0.2857 According to the given proximity rule
        double lonThreshold = 20.0 / 45.0; // 0.4444 According to the given proximity rule

        for (Airport airport : airports) {
            boolean nearLat = Math.abs(flightLat - airport.latitude) <= latThreshold;
            boolean nearLon = Math.abs(flightLon - airport.longitude) <= lonThreshold;

            if (nearLat || nearLon) { // Near either latitude or longitude
                // System.out.println("NEAR " + airport.code + ": " + callSign);
                double dist = Math.sqrt(Math.pow(flightLat - airport.latitude, 2) + Math.pow(flightLon
- airport.longitude, 2)); //calculating the Euclidean distance
                if (nearest == null || dist < minDistance) {

```

```

        nearest = airport;
        minDistance = dist;
    }
}

if (nearest != null) {
    System.out.println("NEAREST " + nearest.city + " (" + nearest.code + "): " + callSign);
    // Airline code is typically the first 3 letters of the call sign
    String airlineCode = callSign.length() >= 3 ? callSign.substring(0, 3).trim() : callSign;
    collector.emit(new Values(nearest.city, nearest.code, airlineCode));
}
} catch (Exception ignored) {}
}

public void declareOutputFields(OutputFieldsDeclarer declarer) {
    declarer.declare(new Fields("airport.city", "airport.code", "call_sign"));
}

public void cleanup() {}

private static class Airport implements Serializable {
    String city, code;
    double latitude, longitude;

    Airport(String city, String code, double latitude, double longitude) {
        this.city = city;
        this.code = code;
        this.latitude = latitude;
        this.longitude = longitude;
    }
}
}

```

AirlineSorter Class:

```

package bolts;

import backtype.storm.task.TopologyContext;
import backtype.storm.topology.BasicOutputCollector;
import backtype.storm.topology.OutputFieldsDeclarer;
import backtype.storm.topology.base.BaseBasicBolt;
import backtype.storm.tuple.Tuple;

import java.util.ArrayList;
import java.util.HashMap;
import java.util.List;
import java.util.Map;

public class AirlineSorter extends BaseBasicBolt {

    // This is used to store the number of flights for each airline at each airport
    private Map<String, Map<String, Integer>> stats;

    // Map from airport code to city name
    private Map<String, String> codeToCity;
}

```



```

@Override
public void prepare(Map stormConf, TopologyContext context) {
    stats = new HashMap<>();
    codeToCity = new HashMap<>();
}

@Override
public void execute(Tuple input, BasicOutputCollector collector) {
    String airportCity = input.getStringByField("airport.city");
    String airportCode = input.getStringByField("airport.code");
    String airlineCode = input.getStringByField("call_sign");

    // Store code -> city mapping
    codeToCity.putIfAbsent(airportCode, airportCity);

    // Keyed by airport code
    stats.putIfAbsent(airportCode, new HashMap<String,Integer>());
    Map<String, Integer> airlineMap = stats.get(airportCode);
    airlineMap.put(airlineCode, airlineMap.getOrDefault(airlineCode, 0) + 1);
}

@Override
public void cleanup() {
    System.out.println("-- Flight Counter --");

    for (Map.Entry<String, Map<String, Integer>> airportEntry : stats.entrySet()) {
        String airport = airportEntry.getKey();
        Map<String, Integer> airlineMap = airportEntry.getValue();
        String cityName = codeToCity.getOrDefault(airport, "Unknown");

        System.out.println("At Airport: " + airport + " (" + cityName + ")");
        int total = 0;

        List<Map.Entry<String, Integer>> sortedEntries = new ArrayList<>(airlineMap.entrySet());
        sortedEntries.sort((a, b) -> b.getValue().compareTo(a.getValue()));

        for (Map.Entry<String, Integer> entry : sortedEntries) {
            System.out.println("\t" + entry.getKey() + ": " + entry.getValue());
            total += entry.getValue();
        }

        System.out.println("\ttotal #flights = " + total);
        System.out.println();
    }
}

@Override
public void declareOutputFields(OutputFieldsDeclarer declarer) {
    // No output fields from final bolt
}
}

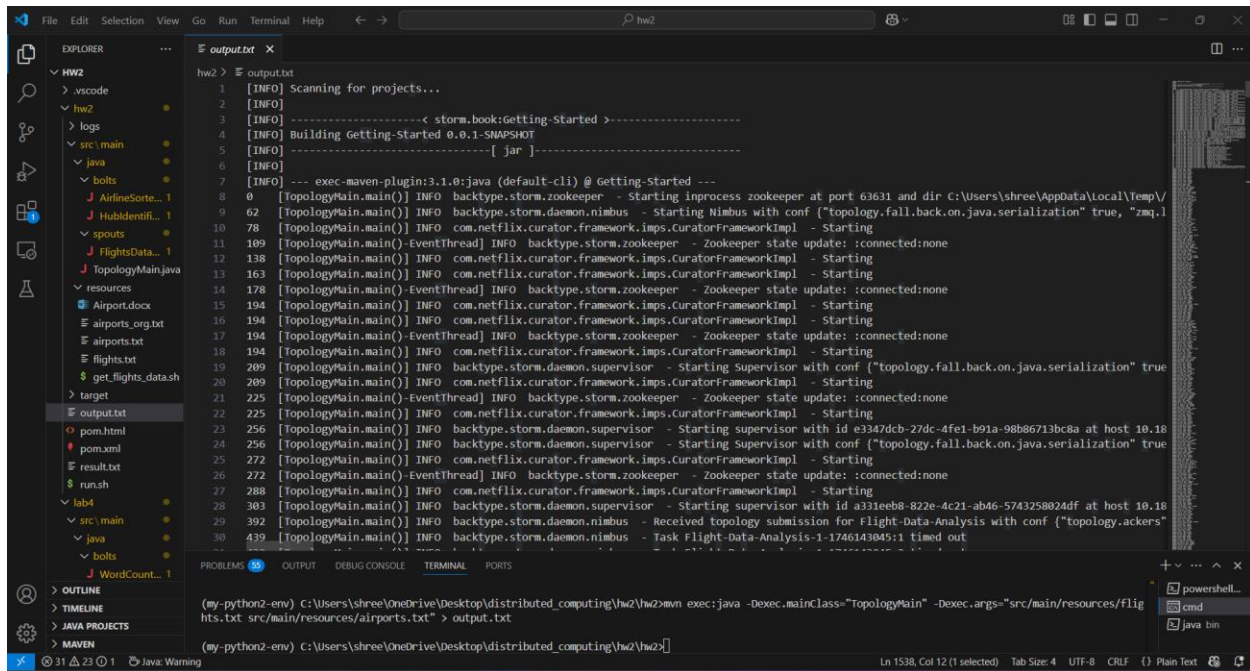
```

## Execution Outputs:

Two scenarios were given to test with the implementation. Below are the outputs for each.

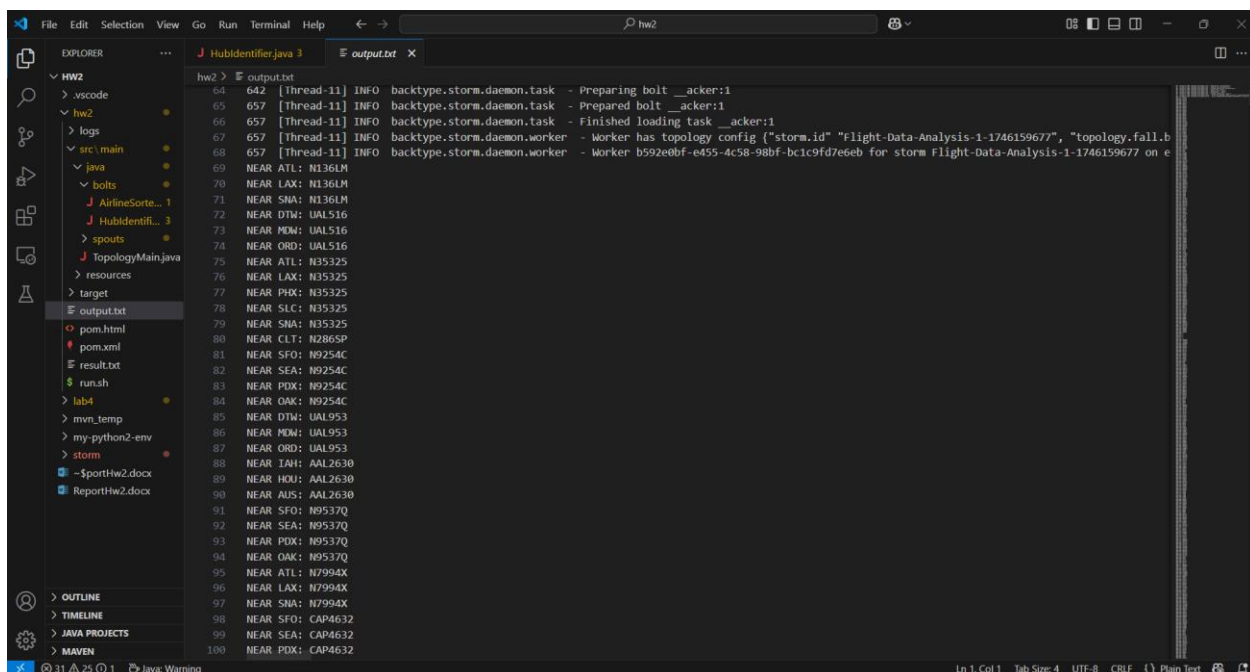
Single-threaded execution:

Execution start.



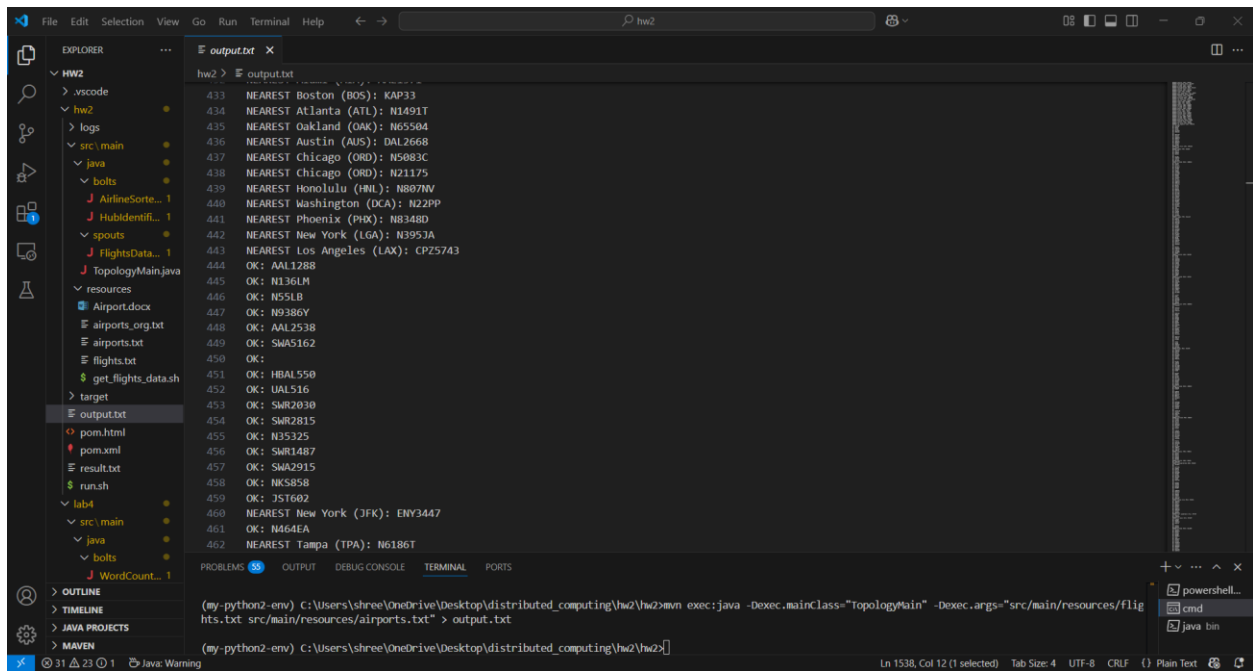
```
1 [INFO] Scanning for projects...
2 [INFO]
3 [INFO] -----< storm.book:Getting-Started >-----
4 [INFO] Building Getting-Started 0.0.1-SNAPSHOT
5 [INFO] [ jar ]-----
6 [INFO] --- exec-maven-plugin:3.1.0:java (default-cli) @ Getting-Started ---
7 [INFO] [TopologyMain.main()] INFO backtype.storm.zookeeper - Starting inprocess zookeeper at port 63631 and dir C:\Users\shree\AppData\Local\Temp\
8 [INFO] [TopologyMain.main()] INFO backtype.storm.daemon.nimbus - Starting Nimbus with conf {"topology.fail.back.on.java.serialization" true, "zmq.l
9 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
10 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
11 [INFO] [TopologyMain.main()-EventThread] INFO backtype.storm.zookeeper - Zookeeper state update: :connected:none
12 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
13 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
14 [INFO] [TopologyMain.main()-EventThread] INFO backtype.storm.zookeeper - Zookeeper state update: :connected:none
15 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
16 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
17 [INFO] [TopologyMain.main()-EventThread] INFO backtype.storm.zookeeper - Zookeeper state update: :connected:none
18 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
19 [INFO] [TopologyMain.main()] INFO backtype.storm.daemon.supervisor - Starting Supervisor with conf {"topology.fail.back.on.java.serialization" true
20 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
21 [INFO] [TopologyMain.main()-EventThread] INFO backtype.storm.zookeeper - Zookeeper state update: :connected:none
22 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
23 [INFO] [TopologyMain.main()] INFO backtype.storm.daemon.supervisor - Starting supervisor with id e3347dc8-27dc-4fe1-b91a-98b6713bc8a at host 10.18
24 [INFO] [TopologyMain.main()] INFO backtype.storm.daemon.supervisor - Starting Supervisor with conf {"topology.fail.back.on.java.serialization" true
25 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
26 [INFO] [TopologyMain.main()-EventThread] INFO backtype.storm.zookeeper - Zookeeper state update: :connected:none
27 [INFO] [TopologyMain.main()] INFO com.netflix.curator.framework.imps.CuratorFrameworkImpl - Starting
28 [INFO] [TopologyMain.main()] INFO backtype.storm.daemon.supervisor - Starting supervisor with id a331eeb8-822e-4c21-ab46-5743258024df at host 10.18
29 [INFO] [TopologyMain.main()] INFO backtype.storm.daemon.nimbus - Received topology submission for Flight-Data-Analysis with conf {"topology.ackers"
30 [INFO] [TopologyMain.main()] INFO backtype.storm.daemon.nimbus - Task Flight-Data-Analysis-1-1746143045:1 timed out
```

One flight mapping multiple airports. Original implementation.



```
64 [Thread-11] INFO backtype.storm.daemon.task - Preparing bolt_acker:1
65 [Thread-11] INFO backtype.storm.daemon.task - Prepared bolt_acker:1
66 [Thread-11] INFO backtype.storm.daemon.task - Finished loading task_acker:1
67 [Thread-11] INFO backtype.storm.daemon.worker - Worker has topology config {"storm.id" "Flight-Data-Analysis-1-1746159677", "topology.fail.b
68 [Thread-11] INFO backtype.storm.daemon.worker - Worker b592e0bf-e455-4c58-98bf-b1c9fd7e6eb for storm Flight-Data-Analysis-1-1746159677 on e
69 NEAR ATL: H136LM
70 NEAR LAX: H136LM
71 NEAR SNA: H136LM
72 NEAR DTW: UAL516
73 NEAR MDW: UAL516
74 NEAR ORD: UAL516
75 NEAR ATL: R35325
76 NEAR LAX: R35325
77 NEAR PDX: R35325
78 NEAR SLC: R35325
79 NEAR SNA: R35325
80 NEAR CLT: R2865P
81 NEAR SFO: R9254C
82 NEAR SEA: R9254C
83 NEAR PDX: R9254C
84 NEAR OAK: R9254C
85 NEAR DTW: UAL953
86 NEAR MDW: UAL953
87 NEAR ORD: UAL953
88 NEAR IAH: AAL2630
89 NEAR HOU: AAL2630
90 NEAR AUS: AAL2630
91 NEAR SFO: R9537Q
92 NEAR SEA: R9537Q
93 NEAR PDX: R9537Q
94 NEAR OAK: R9537Q
95 NEAR ATL: R7994X
96 NEAR LAX: R7994X
97 NEAR SNA: R7994X
98 NEAR SFO: CAP4632
99 NEAR SEA: CAP4632
100 NEAR PDX: CAP4632
```

## Receiving acknowledgement



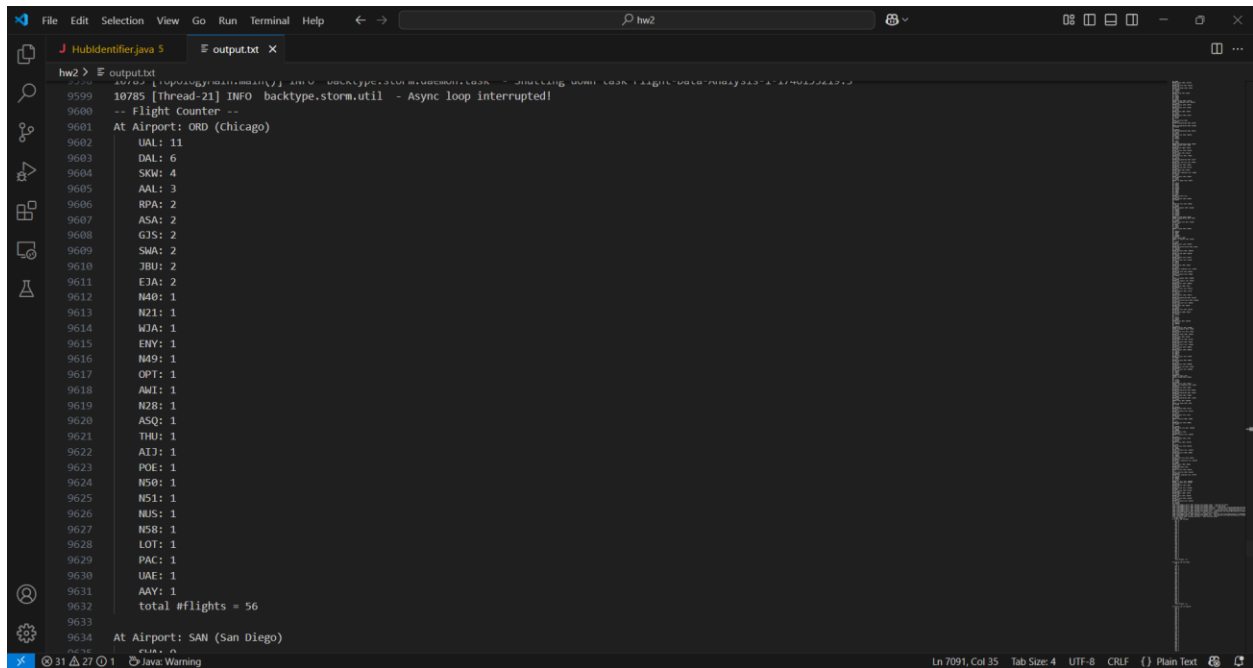
```
433 NEAREST Boston (BOS): KAP33
434 NEAREST Atlanta (ATL): N1491T
435 NEAREST Oakland (OAK): N65504
436 NEAREST Austin (AUS): DAL2668
437 NEAREST Chicago (ORD): N5083C
438 NEAREST Chicago (ORD): N21175
439 NEAREST Honolulu (HNL): N807NV
440 NEAREST Washington (DCA): N229P
441 NEAREST Phoenix (PHX): N834BD
442 NEAREST New York (LGA): N395JA
443 NEAREST Los Angeles (LAX): CP25743
444 OK: AAL1288
445 OK: N136LM
446 OK: N55LB
447 OK: N9386Y
448 OK: AAL2538
449 OK: SWA5162
450 OK:
451 OK: HBAL550
452 OK: UAL516
453 OK: SWR2030
454 OK: SWR2815
455 OK: N35325
456 OK: SWR1487
457 OK: SWA2915
458 OK: NKS858
459 OK: JST602
460 NEAREST New York (JFK): ENY3447
461 OK: N464EA
462 NEAREST Tampa (TPA): N6186T
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

(my-python2-env) C:\Users\shree\OneDrive\Desktop\distributed\_computing\hw2\mvn exec:java -Dexec.mainClass="TopologyMain" -Dexec.args="src/main/resources/flights.txt src/main/resources/airports.txt" > output.txt

(my-python2-env) C:\Users\shree\OneDrive\Desktop\distributed\_computing\hw2\hw2

## Original Implementation- Mapping near airport.



```
9599 10785 [Thread-21] INFO backtype.storm.util - Async loop interrupted!
9600 -- Flight counter --
9601 At Airport: ORD (Chicago)
9602 UAL: 11
9603 DAL: 6
9604 SKW: 4
9605 AAL: 3
9606 RPA: 2
9607 ASA: 2
9608 GJS: 2
9609 SWA: 2
9610 JBU: 2
9611 EJA: 2
9612 N40: 1
9613 N21: 1
9614 WJA: 1
9615 ENY: 1
9616 N49: 1
9617 OPT: 1
9618 AMT: 1
9619 N28: 1
9620 ASQ: 1
9621 THU: 1
9622 ATJ: 1
9623 POE: 1
9624 NS0: 1
9625 NS1: 1
9626 NUS: 1
9627 NS8: 1
9628 LOT: 1
9629 PAC: 1
9630 UAL: 1
9631 AAY: 1
9632 total #flights = 56
9633
9634 At Airport: SAN (San Diego)
```

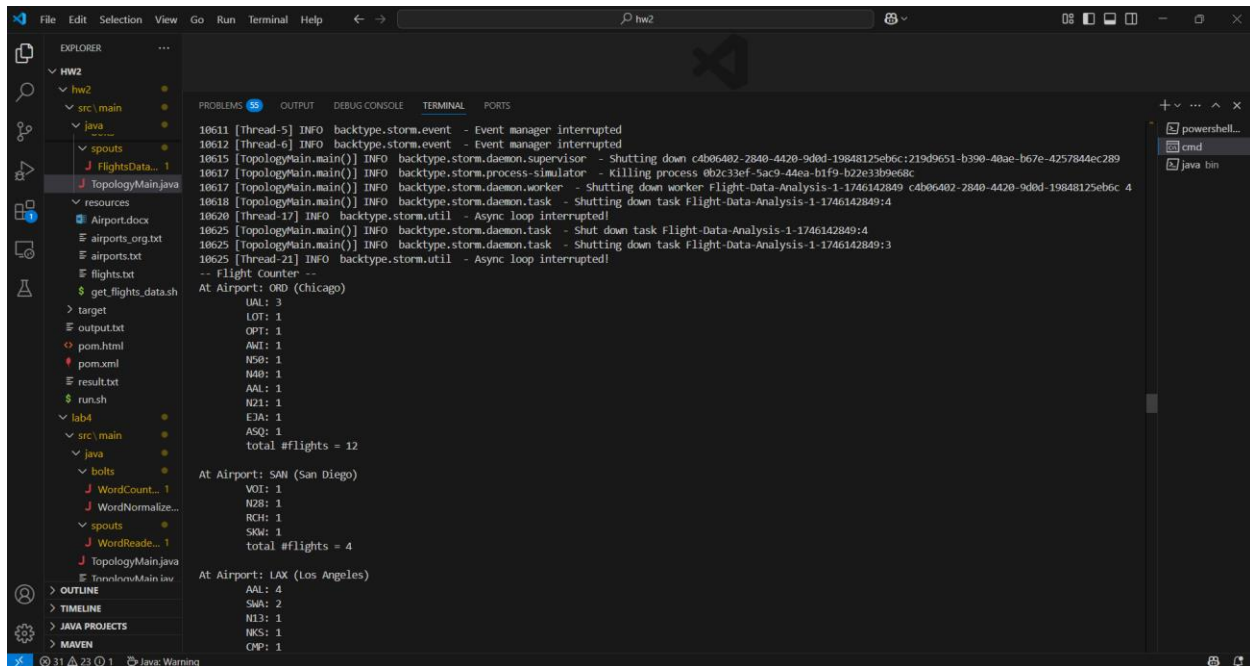
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

(my-python2-env) C:\Users\shree\OneDrive\Desktop\distributed\_computing\hw2\mvn exec:java -Dexec.mainClass="TopologyMain" -Dexec.args="src/main/resources/flights.txt src/main/resources/airports.txt" > output.txt

(my-python2-env) C:\Users\shree\OneDrive\Desktop\distributed\_computing\hw2\hw2

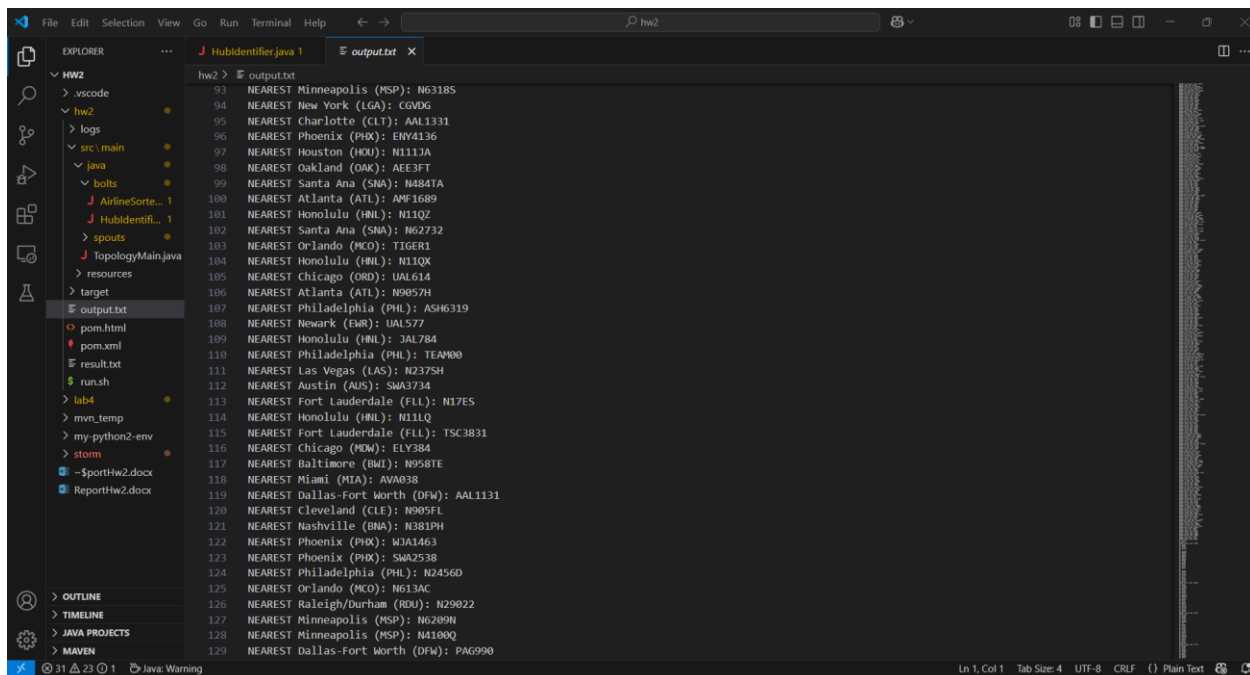
Additional feature(s):

Sorted and filtered output.



```
10611 [Thread-5] INFO backtype.storm.event - Event manager interrupted
10612 [Thread-6] INFO backtype.storm.event - Event manager interrupted
10613 [TopologyMain.main()] INFO backtype.storm.daemon.supervisor - Shutting down c4b06402-2840-4420-9d0d-19848125eb6c:219d9651-b390-40ae-b67e-4257844ec289
10617 [TopologyMain.main()] INFO backtype.storm.process.simulator - Killing process 0b2c33ef-5ac9-44ea-b1f9-b22e33b9e68c
10617 [TopologyMain.main()] INFO backtype.storm.daemon.worker - Shutting down worker Flight-Data-Analysis-1-1746142849:4
10618 [TopologyMain.main()] INFO backtype.storm.daemon.task - Shutting down task Flight-Data-Analysis-1-1746142849:4
10620 [Thread-17] INFO backtype.storm.util - Async loop interrupted!
10625 [TopologyMain.main()] INFO backtype.storm.daemon.task - Shut down task Flight-Data-Analysis-1-1746142849:4
10625 [TopologyMain.main()] INFO backtype.storm.daemon.task - Shutting down task Flight-Data-Analysis-1-1746142849:3
10625 [Thread-21] INFO backtype.storm.util - Async loop interrupted!
-- Flight Counter --
At Airport: ORD (Chicago)
    UAL: 3
    LOT: 1
    OPT: 1
    AMI: 1
    HSR: 1
    NAR: 1
    AAL: 1
    N21: 1
    EJA: 1
    ASQ: 1
    total #flights = 12
At Airport: SAN (San Diego)
    VOI: 1
    N28: 1
    RCH: 1
    SMD: 1
    total #flights = 4
At Airport: LAX (Los Angeles)
    AAL: 4
    SWA: 2
    N13: 1
    BKS: 1
    CXP: 1
```

Mapping to only one nearest airport



```
93 NEAREST Minneapolis (MSP): N6318S
94 NEAREST New York (LGA): CGV0G
95 NEAREST Charlotte (CLT): AAL1331
96 NEAREST Phoenix (PHX): ENY4136
97 NEAREST Houston (HOU): N111JA
98 NEAREST Oakland (OAK): AEE3FT
99 NEAREST Santa Ana (SNA): N484TA
100 NEAREST Atlanta (ATL): AMF1689
101 NEAREST Honolulu (HNL): N11QZ
102 NEAREST Santa Ana (SNA): N62732
103 NEAREST Orlando (MCO): TIGER1
104 NEAREST Honolulu (HNL): N11QX
105 NEAREST Chicago (ORD): UAL614
106 NEAREST Atlanta (ATL): N9057H
107 NEAREST Philadelphia (PHL): ASH6319
108 NEAREST Newark (EWR): UAL577
109 NEAREST Honolulu (HNL): JAL784
110 NEAREST Philadelphia (PHL): TEA400
111 NEAREST Las Vegas (LAS): N237SH
112 NEAREST Austin (AUS): SWA3734
113 NEAREST Fort Lauderdale (FLL): N17ES
114 NEAREST Honolulu (HNL): N11LQ
115 NEAREST Fort Lauderdale (FLL): TSC3831
116 NEAREST Chicago (MDW): ELY384
117 NEAREST Baltimore (BWI): N958TE
118 NEAREST Miami (MIA): AWA038
119 NEAREST Dallas-Fort Worth (DFW): AAL1131
120 NEAREST Cleveland (CLE): N905FL
121 NEAREST Nashville (BNA): N381PH
122 NEAREST Phoenix (PHX): N3A1463
123 NEAREST Phoenix (PHX): SWA2538
124 NEAREST Philadelphia (PHL): N2456D
125 NEAREST Orlando (MCO): N613AC
126 NEAREST Raleigh/Durham (RDU): N29022
127 NEAREST Minneapolis (MSP): N6209N
128 NEAREST Minneapolis (MSP): N4100Q
129 NEAREST Dallas-Fort Worth (DFW): PAG990
```

## Discussions:

### Additional Features:

1. Sorting the airline based on the number of flights it has in an airport.
2. More accurate flight information:
  - a. Filtered the airplanes according to different parameters such as altitude, vertical rate.
  - b. Improved the airplanes mappings to its nearest airport

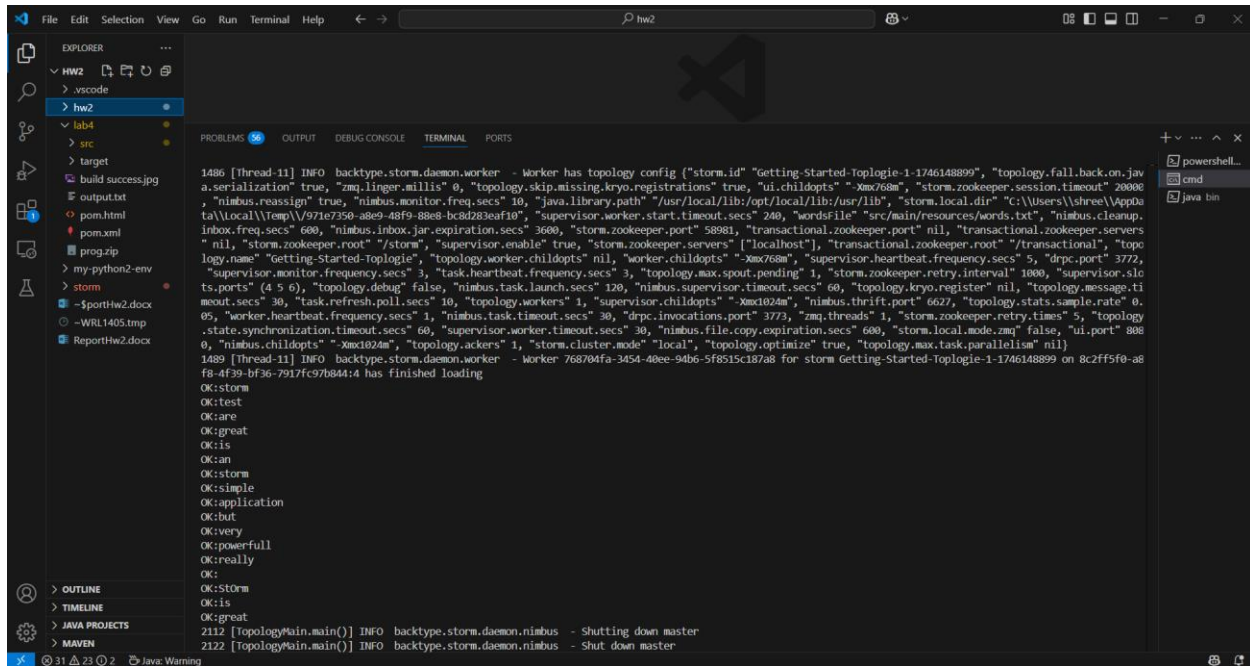
To implement the feature of sorting airlines based on how many flights they have at each airport, the `AirlineSorter` class uses a nested map structure called `stats`. This map keeps track of flight counts, with the airport code as the outer key and a map of airline codes and their respective counts as the inner values. The airline code is extracted as the first three characters of the call sign, which is passed from the `HubIdentifier` class. As flight data comes in during the `execute()` method, the corresponding airline's count is updated. Then, in the `cleanup()` method, each airport's airline data is collected, turned into a list, and sorted in descending order based on flight count. Finally, the class prints a summary for each airport showing its city, the airlines that operated there, and how many flights each had.

For the Implementation of more accurate flight data. The following implementations are done.

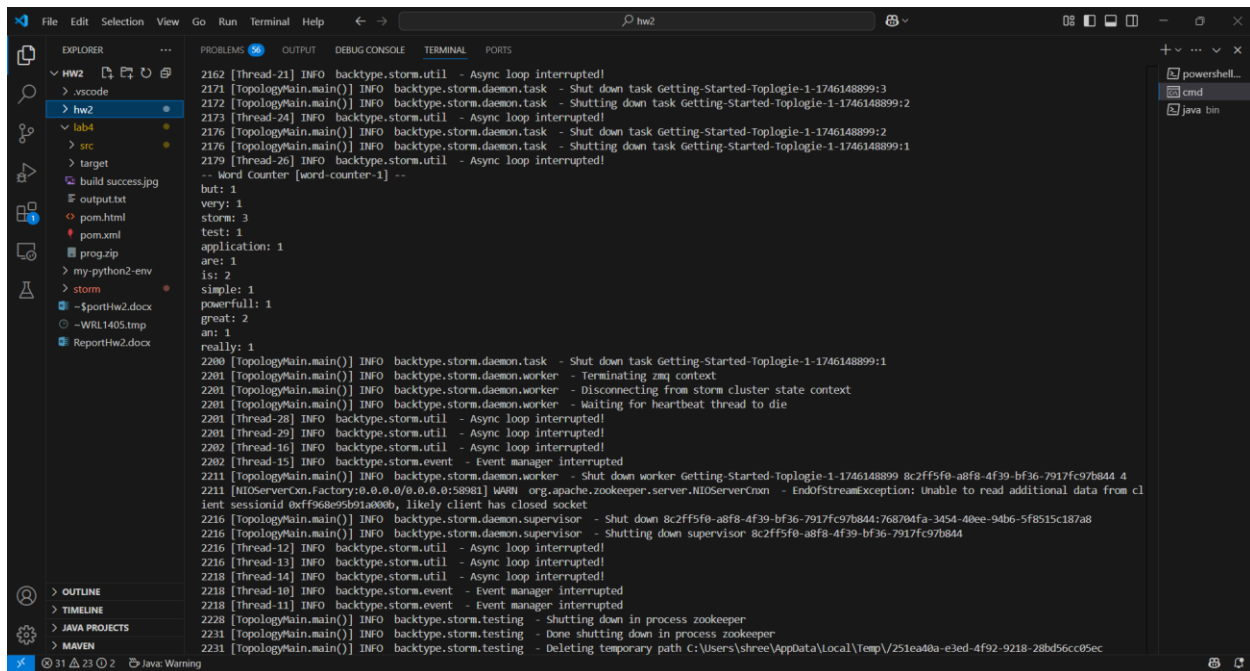
In the `HubIdentifier` class, a set of filters is used to focus on flights that are actually relating with an airport, such as those taking off, landing, or on the ground, rather than those simply flying overhead. The class first checks if the flight has a valid call sign, as missing call signs often indicate incomplete or less useful data. It then filters out flights that are airborne but not changing altitude, which usually means they are just cruising by. Flights moving faster than 200 meters per second are also excluded, since they are likely passing over the area at high speed. Additionally, flights above 1000 feet are not considered, as the focus is on aircraft closer to the ground or near the airport. Once these filters are applied, the remaining flights are checked for proximity to known airports based on whether they are within about 20 miles in either latitude or longitude. This approach ensures that only flights that are related or connected to an airport are included in the topology.

In the initial sample output, it was observed that some flights were being mapped to multiple airports due to the broad proximity rule, which considers a flight near any airport if it is within 20 miles in either latitude or longitude. While this approach captures potential proximity, it can also introduce ambiguity and inflate the data, since a single flight might be linked to several airports. To avoid this confusion, the `HubIdentifier` class was designed to identify only the nearest airport for each flight. Instead of emitting a flight for every matching airport within the proximity range, the implementation keeps track of the closest one by calculating the Euclidean distance between the flight's coordinates and each airport. This is done by looping through all airports and updating the minimum distance and corresponding airport whenever a closer one is found. As a result, each flight is uniquely and more accurately associated with just one airport, which simplifies the analysis and leads to more meaningful flight-to-airport mappings.

## Lab 4:



```
1486 [Thread-11] INFO backtype.storm.daemon.worker - Worker has topology config {"storm.id": "Getting-Started-Topologie-1-1746148899", "topology.fail.back.on.java.serialization": true, "zmq.linger.millis": 0, "topology.skip.missing.kryo.registrations": true, "ui.childopts": "-Xmx768m", "storm.zookeeper.session.timeout": 20000, "nimbus.monitor.freq.secs": 10, "nimbus.monitor.freq.secs": 10, "java.library.path": "/usr/local/lib:/opt/local/lib:/usr/lib", "storm.local.dir": "C:\\Users\\shree\\AppData\\Local\\Temp\\V971e7350-a8e9-48f9-88e8-bcd283eaf10", "supervisor.worker.start.timeout.secs": 240, "wordsFile": "src/main/resources/words.txt", "nimbus.cleanup.inbox.freq.secs": 600, "nimbus.inbox.jar.expiration.secs": 3600, "storm.zookeeper.port": 58981, "transactional.zookeeper.port": nil, "transactional.zookeeper.servers": nil, "storm.zookeeper.root": "/storm", "supervisor.enable": true, "storm.zookeeper.servers": ["localhost"], "transactional.zookeeper.root": "/transactional", "topology.name": "Getting-Started-Topologie", "topology.worker.childopts": nil, "worker.childopts": "-Xmx768m", "supervisor.heartbeat.frequency.secs": 5, "drpc.port": 3772, "supervisor.monitor.frequency.secs": 3, "task.heartbeat.frequency.secs": 3, "topology.max.spout.pending": 1, "storm.zookeeper.retry.interval": 1800, "supervisor.slave.ports": (4 5 6), "topology.debug": false, "nimbus.task.launch.secs": 120, "nimbus.supervisor.timeout.secs": 60, "topology.kryo.register": nil, "topology.message.timeout.secs": 30, "task.refresh.poll.secs": 10, "topology.workers": 1, "supervisor.childopts": "-Xmx1024m", "nimbus.thrift.port": 6627, "topology.stats.sample.rate": 0.05, "worker.heartbeat.frequency.secs": 1, "nimbus.task.timeout.secs": 30, "drpc.invocations.port": 3773, "zmq.threads": 1, "storm.zookeeper.retry.times": 5, "topology.state.synchronization.timeout.secs": 60, "supervisor.worker.timeout.secs": 30, "nimbus.file.copy.expiration.secs": 600, "storm.local.mode.zmq": false, "ui.port": 8080, "nimbus.childopts": "-Xmx1024m", "topology.ackers": 1, "storm.cluster.mode": "local", "topology.optimize": true, "topology.max.task.parallelism": nil}
1489 [Thread-11] INFO backtype.storm.daemon.worker - Worker 768704fa-3454-40ee-94b6-5f8515c187a8 for storm Getting-Started-Topologie-1-1746148899 on 8c2ff5f0-a8f8-4f39-bf36-7917fc97b844:4 has finished loading
OK:storm
OK:test
OK:care
OK:great
OK:is
OK:an
OK:storm
OK:simple
OK:application
OK:but
OK:very
OK:powerfull
OK:really
OK:
OK:Storm
OK:is
OK:great
2112 [TopologyMain.main()] INFO backtype.storm.daemon.nimbus - Shutting down master
2122 [TopologyMain.main()] INFO backtype.storm.daemon.nimbus - Shut down master
```



```
2162 [Thread-21] INFO backtype.storm.util - Async loop interrupted!
2171 [TopologyMain.main()] INFO backtype.storm.daemon.task - Shut down task Getting-Started-Topologie-1-1746148899:3
2172 [TopologyMain.main()] INFO backtype.storm.daemon.task - Shutting down task Getting-Started-Topologie-1-1746148899:2
2173 [Thread-24] INFO backtype.storm.util - Async loop interrupted!
2176 [TopologyMain.main()] INFO backtype.storm.daemon.task - Shut down task Getting-Started-Topologie-1-1746148899:2
2176 [TopologyMain.main()] INFO backtype.storm.daemon.task - Shutting down task Getting-Started-Topologie-1-1746148899:1
2179 [Thread-26] INFO backtype.storm.util - Async loop interrupted!
-- Word Counter [word-counter-1] --
but: 1
very: 1
storm: 3
test: 1
application: 1
are: 1
is: 2
simple: 1
powerfull: 1
great: 2
an: 1
really: 1
2200 [TopologyMain.main()] INFO backtype.storm.daemon.task - Shut down task Getting-Started-Topologie-1-1746148899:1
2201 [TopologyMain.main()] INFO backtype.storm.daemon.worker - Terminating zmq context
2201 [TopologyMain.main()] INFO backtype.storm.daemon.worker - Disconnecting from storm cluster state context
2201 [Thread-28] INFO backtype.storm.util - Async loop interrupted!
2201 [Thread-29] INFO backtype.storm.util - Async loop interrupted!
2202 [Thread-16] INFO backtype.storm.util - Async loop interrupted!
2202 [Thread-15] INFO backtype.storm.event - Event manager interrupted
2211 [TopologyMain.main()] INFO backtype.storm.daemon.worker - Shut down worker Getting-Started-Topologie-1-1746148899 8c2ff5f0-a8f8-4f39-bf36-7917fc97b844 4
2211 [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:58981] WARN org.apache.zookeeper.server.NIOServerCxn - EndOfStreamException: Unable to read additional data from client sessionid 0xf968e95b91a000b, likely client has closed socket
2216 [TopologyMain.main()] INFO backtype.storm.daemon.supervisor - Shut down 8c2ff5f0-a8f8-4f39-bf36-7917fc97b844:768704fa-3454-40ee-94b6-5f8515c187a8
2216 [TopologyMain.main()] INFO backtype.storm.daemon.supervisor - Shutting down supervisor 8c2ff5f0-a8f8-4f39-bf36-7917fc97b844
2216 [Thread-12] INFO backtype.storm.util - Async loop interrupted!
2218 [Thread-14] INFO backtype.storm.util - Async loop interrupted!
2218 [Thread-10] INFO backtype.storm.event - Event manager interrupted
2218 [Thread-11] INFO backtype.storm.event - Event manager interrupted
2228 [TopologyMain.main()] INFO backtype.storm.testing - Shutting down in process zookeeper
2231 [TopologyMain.main()] INFO backtype.storm.testing - Done shutting down in process zookeeper
2231 [TopologyMain.main()] INFO backtype.storm.testing - Deleting temporary path C:\\Users\\shree\\AppData\\Local\\Temp\\V251ea40a-e3ed-4f92-9218-28bd56cc05ec
```